

Analyzing Goodput in Different Network Scenarios: Ethernet vs Wi-Fi

Conforto Francesco Paolo, D'Amico Dennis, Drago Giovanni, Lo Tufo Stefano

1 GOALS OF THE LABORATORY

This study focuses on evaluating the performance characteristics of different network scenarios, ranging from wired (Ethernet) to wireless (WiFi) connections, through the application of TCP and UDP protocols. The experiment aims to study how these protocols respond to different network scenarios. By systematically analyzing their behavior, we seek to quantify their efficiency, reliability, and overall effectiveness in data transmission.

For each test, we will compute the theoretical Goodput, which quantifies the amount of useful application-layer data successfully transmitted over the network, excluding interference/congestion, retransmissions, and control messages. This theoretical value will be compared against the measured Goodput obtained from practical experiments to assess discrepancies between ideal and real-world performance.

2 THEORY & EXPECTED RESULT

In network communication, understanding the efficiency and effectiveness of data transmission is essential for optimizing performance.

Two critical metrics in this context are Efficiency and Goodput, together, these metrics provide a comprehensive view of both the actual data delivery and the overall performance of the network. This section explores these two concepts, providing their definitions and formulas.

2.1 Efficiency

Transmission efficiency (η) is a measure of the proportion of useful application-layer data successfully transmitted relative to the total amount of data sent. It quantifies how efficiently the available bandwidth is utilized for delivering actual payload data.

Follows the formula used to calculate the efficiency:

$$\eta = \frac{Data_{L7}}{Data_{L7} + Overhead_{L4} + Overhead_{L3} + Overhead_{L2}} \quad (1)$$

- $Data_{L7}$ refers to the application layer data (Layer 7 of the OSI model). It is the useful data that is passed from the application (such as a file, a message, or any other form of application data) that needs to be transmitted across the network.
- $Overhead_{L4}$ refers to the transport layer overhead (Layer 4 in the OSI model), which includes the headers and control information for protocols like TCP or UDP.
- $Overhead_{L3}$ refers to the network layer overhead (Layer 3, OSI model) provided by the IP header. This header, which includes source/destination IPs, TTL, protocol, and checksum, is essential for routing packets.

- $Overhead_{L2}$ refers to the Data Link Layer overhead (Layer 2 in the OSI model), which includes the frame header and trailer for both Ethernet and WiFi technologies.

2.2 Goodput

Goodput is a key performance metric that measures the rate of useful data delivered to the application layer, excluding retransmissions, overhead, and control messages. Follows the formula to calculate the Goodput:

$$G = \frac{D}{T} \quad (2)$$

- D is the total amount of successfully received application-layer data (bits).
- T is the total transmission time (seconds).

A different way to express Goodput takes into account the efficiency (η) of the transmission and the channel capacity (C). This leads to the formula:

$$G = \eta \cdot C \quad (3)$$

- η is the efficiency of the transmission
- C is the channel capacity (i.e., the maximum possible transmission rate, typically given in bits per second).

2.3 Expected Results

Theoretical Goodput (T.G.) differs from practical Goodput due to various real-world factors that impact network performance. In theory, Goodput is calculated based on ideal conditions, in practice, issues such as network congestion, packet loss, retransmissions, and network interference can all reduce the actual Goodput. These factors introduce delays, reduce efficiency, and result in a lower practical Goodput than the theoretical value.

3 TOOLS

In this section we briefly see the commands used to perform the tests.

3.1 Iperf3

Iperf3 is a tool that, by generating traffic, estimates some key metrics to calculate the real average Goodput. By default, it records data every second on ten samples, then calculates the average Bitrate for both the sender and the receiver. The average Bitrate of the receiver is used as the real Goodput, as defined in the formula in the previous section.

3.2 The commands

`iperf3 -c <server_ip> -J -t 10`

With this command we can test the TCP connection between two computers for the duration of 10 seconds and output the result in JSON format.

iperf3 -c <server_ip> -J -t 10 -u -b <value> -l 1472

With this command we can test the UDP connection between two computers for the duration of 10 seconds and also it permits to choose the bitrate value thanks to the -b flag. The -l flag permits to set the fragmentation size.

3.3 Wireshark

Wireshark is a widely used network protocol analyzer that captures and inspects data packets traveling through a network. In performance testing, it monitors protocol behavior under various conditions and loads.

Additionally, Wireshark enables the creation of custom graphs based on specific network parameters. For our tests, we used it to plot sequence numbers, throughput, and bitrate over time.

4 SCENARIOS AND RESULTS

This section provides a detailed analysis of three scenarios and the corresponding results. To perform the test we used the Script 1 attached in the appendix 1

4.1 Ethernet to Ethernet

For the first case of study we connected two Windows computers through an Ethernet cable with capacity **1 Gbps** as shown in the Figure 1. From this configuration, we expect a goodput of **949 Mbps** for TCP and **957 Mbps** for UDP as calculated in the formula (3), assuming that the efficiency changes according to the protocol. In the TCP case, the efficiency (η) is equal to 0,949 while in UDP is 0,957.




Figure 1: Ethernet to Ethernet configuration

Follows the table with the results obtained from the Iperf3 tests:

Protocol	T.G.	AVG	Max	Min	Dev Std
TCP	949	947.10	948.00	944.00	1.29
UDP t=5	957	926.00	934.00	915.00	9.13
UDP t=10	957	876.43	910.00	855.00	16.8

Table 1: Results of throughput measurements in Mbps (ETH to ETH)

In the TCP scenario, the measured performance aligns closely with theoretical expectations, with an average bitrate of 947.10 Mbps compared to the theoretical 949 Mbps. In the UDP tests, two distinct behaviors were observed: one using test with a duration of 5 seconds and another one with a duration of 10 seconds. In the first test case, the transmission rate averaged around 954 Mbps; however, a packet loss of approximately 2.45% reduced the effective Goodput, likely due to receiver buffer overflow. In the second test case, instead, a notable issue occurred: after about seven seconds, the transmission rate dropped sharply from roughly 954 Mbps to between 750 and 800 Mbps, suggesting potential factors such as

buffer overflows or misconfigurations.

This strange behavior is not present when a Linux and a Windows machine are connected to each other, where the average slightly exceeds 950 Mbps, always remaining below the theoretical Goodput.

In the following section, graphs will be presented to visually illustrate these differences and provide a clearer understanding of the behavior observed during the tests.

4.2 Wi-Fi to Wi-Fi

For the second case of study, we connected two Windows computers, as shown in Figure 2, through a Wi-Fi connection with a capacity of **400 Mbps**.




Figure 2: Wi-Fi to Wi-Fi configuration

From this configuration, we expect a goodput calculated as follows:

$$\text{Goodput} = 400 \times \eta_{\text{Wi-Fi}} \times \eta_{\text{TCP/IP or UDP/IP}} \times 0.5 \quad (4)$$

- **400 Mbps Factor:** This value represents the theoretical maximum capacity of the Wi-Fi link used in our experiment.
- $\eta_{\text{Wi-Fi}}$: We denote by Wi-Fi Efficiency the efficiency of data transmission at the MAC level of the **802.11ac** protocol, which in our case is equal to **0.8**.
- $\eta_{\text{TCP/IP or UDP/IP}}$: - In the case of **TCP/IP**, the Transport and Network Protocol Efficiency is calculated based on the formula (1):

$$\eta_{\text{TCP/IP}} = \frac{1460}{1460 + 20 + 20} = 0.9733 \quad (5)$$

In the case of **UDP/IP**, the efficiency is slightly higher and is calculated following the expression (1) as:

$$\eta_{\text{UDP/IP}} = \frac{1460}{1460 + 8 + 20} = 0.9813 \quad (6)$$

- **0.5 Factor:** Each packet must be transmitted twice, once from the sender to the Access Point, and again from the Access Point to the receiver

Based on these values, we estimate a goodput of **155.73 Mbps** for TCP and **157.01 Mbps** for UDP.

Follows the table with the results obtained from the Iperf3 tests:

Protocol	T.G.	AVG	Max	Min	Dev Std
TCP	155.73	123.57	129.57	116.74	4.40
UDP	157.01	141.38	150.16	127.89	7.84

Table 2: Results of throughput measurements in Mbps (Wi-Fi to Wi-Fi)

In both TCP and UDP scenarios, the measured goodput is slightly lower than the theoretical values. For TCP, the average measured

bitrate was 123.57 Mbps versus a theoretical 155.73 Mbps. In the UDP scenario, the measured goodput was 141.38 Mbps compared to 157.01 Mbps theoretically. This reduction is likely due to packet loss from transient network conditions, such as congestion or interference, since the tests were performed on the university network.

4.3 Wi-Fi to Ethernet

For the third case of study we used two Windows computers. The sender was connected with the AP through an ethernet cable with capacity of **1 Gbps** while the receiver was connected through Wi-Fi with the AP as shown in the Figure 3.

From this configuration, we expect a goodput calculated as follows:

$$\text{Goodput} = 866.7 \times \eta_{\text{Wi-Fi}} \times \eta_{\text{TCP/IP or UDP/IP}} \quad (7)$$

In this configuration, we consider the Wi-Fi Efficiency and the Transport and Network Protocol Efficiency to be the same as in the Wi-Fi to Wi-Fi configuration ($\eta_{\text{Wi-Fi}}=0.8$, $\eta_{\text{TCP/IP or UDP/IP}}=0.9733$ or 0.9813). This choice was made because, given the network topology, the bottleneck is located at the device connected via Wi-Fi. From this configuration, we expect a goodput of **674.87** for TCP and **680.42** for UDP.




Figure 3: Wi-Fi to Ethernet configuration

Follows the table with the results obtained from the Iperf3 tests:

Protocol	T.G.	AVG	Max	Min	Dev Std
TCP	674.87	604.00	633.00	584.00	16.6
UDP	680.42	653.30	670.00	629.00	14.4

Table 3: Results of throughput measurements in Mbps (Wi-Fi to Ethernet)

As shown, the measured TCP Goodput is consistently lower than the Theoretical Goodput (T.G.), likely due to the limited capacity of the wireless channel, leading to frequent retransmissions and packet loss. This congestion could explain the gap between the measured values and the theoretical prediction.

In the UDP scenario, the results are more aligned with the Theoretical Goodput, likely due to the absence of retransmissions, allowing for a much higher transmission speed. However, this comes at the cost of reliability, as an average packet loss of 30% was observed. This loss occurs because packets are transmitted at the maximum capacity of the Ethernet link before transitioning to a wireless channel, which, being inherently less efficient, cannot sustain the same performance. caption

5 GRAPHIC COMPARISON

This section compares the most significant metrics obtained using Wireshark. The graphs show two different parameters: Throughput

and Bitrate. All other related graphs are provided in the Appendix (A.2, A.3).

5.1 Throughput

This section presents throughput graphs for three network scenarios: Ethernet-to-Ethernet (4), Wi-Fi-to-Wi-Fi (6), and Ethernet-to-Wi-Fi (5).




Figure 4: Ethernet to Ethernet throughput for TCP




Figure 5: Ethernet to Wi-Fi throughput for TCP





Figure 6: Wi-Fi to Wi-Fi throughput for TCP


- **Ethernet-to-Ethernet:** Throughput is generally high and stable due to the wired connection, which offers high capacity and minimal interference. (4)
- **Ethernet-to-Wi-Fi:** Throughput is noticeably lower compared to the wired scenario. This reduction is likely due to differences in channel capacity and efficiency, as well as interference and congestion that may cause TCP retransmissions. (5)
- **Wi-Fi-to-Wi-Fi:** Throughput measurements tend to be closer to the theoretical values, despite the inherent variability of wireless channels. (6)

349 5.2 Bitrate

350 This section presents bitrate graph comparisons for TCP and UDP
 351 in three network scenarios.



352 **Figure 7:** Ethernet to Ethernet bits per sec for TCP




353 **Figure 8:** Wi-Fi to Wi-Fi bits per sec for TCP

354 The performance analysis of Ethernet-Ethernet (7) and Wi-Fi-Wi-Fi
 355 (8) in TCP transmission, reveals significant differences in stability
 356 and bitrate fluctuations.

357 Ethernet shows a stable trend with minimal variability, thanks to
 358 the reliable wired connection, low latency, and absence of interference.
 359 This results in consistent throughput.

360 In contrast, Wi-Fi experiences noticeable bitrate peaks and drops
 361 due to factors like wireless interference, congestion, and signal
 362 attenuation, leading to less stable throughput.


363 These differences confirm that Ethernet offers more predictable
 364 performance, while Wi-Fi, although more flexible, is more sensitive
 365 to fluctuations that can impact sensitive applications.



366 **Figure 9:** Ethernet to Ethernet bits per sec for UDP, 5 seconds case
 367 Windows-Windows



398 **Figure 10:** Ethernet to Ethernet bits per sec for UDP, 10 seconds case
 399 Windows-Windows



400 **Figure 11:** Ethernet to Ethernet bits per sec for UDP, 10 seconds case
 401 Ubuntu-Windows

402 The UDP performance was analyzed in three scenarios with different OS configurations and detection intervals, revealing key
 403 differences in stability and throughput.

- 404 • **Scenario 1:** Two Windows systems with a 5-second test show stable throughput. (9)
- 405 • **Scenario 2:** Two Windows systems with a 10-second test experience a drop to 750-800 Mbps, likely due to buffering or congestion management, possibly influenced by the use of the APIPA protocol for IP assignment. (10)
- 406 • **Scenario 3:** A Linux and a Windows system maintain throughput close to the expected goodput even with a 10-second test. In this case, IPs are assigned statically, leading to more efficient UDP handling on Linux. (11)

407 These results highlight that OS configuration, measurement interval, and APIPA usage in Windows-only setups significantly impact
 408 UDP performance. Linux systems mitigate bitrate drops seen in
 409 Windows-only configurations, likely due to differences in packet
 410 handling.

411 6 CONCLUSION

412 From the experiments that we have conducted we obtained interesting results to be analyzed:

- 413 • For Ethernet to Ethernet we had a stable bitrate for TCP while for UDP we experienced a significant drop of the transmission rate in the ten seconds test.
- 414 • For Wi-Fi to Wi-Fi, we experienced a result more aligned with the theoretical one. The little discrepancy is probably caused by the network congestion or interferences, since the tests were performed into the university network.
- 415 • For the Ethernet to WiFi configuration, there is a significant discrepancy in TCP between the theoretical Goodput and the measured value. In contrast, UDP measurements are much closer to the theoretical value.

416 Overall, while some results aligned well with theoretical expectations, others deviated due to various influencing factors. These findings highlight the impact of network conditions, protocol characteristics, and transmission environments on performance, underscoring the need for further investigations to refine theoretical models and optimize network performance in heterogeneous environments.

465

A APPENDIX

466

A.1 Code

467

Follows the Python code used to perform the tests.

470

471

472

473

474

475

476

477

478

479

480

481

482

483

484

485

486

487

488

489

490

491

492

493

494

495

496

497

498

499

500

501

502

503

504

505

506

507

508

509

510

511

512

513

514

515

516

517

518

519

520

521

522

523

524

525

526

527

528

529

530

531

532

533

534

535

536

537

538

539

540

541

542

543

544

545

546

547

548

549

550

551

552

553

554

555

556

557

558

559

560

561

562

563

564

565

566

567

568

569

570

571

572

573

574

575

576

577

578

579

580

581

582

583

584

585

586

587

588

589

590

591

592

593

594

595

596

597

598

599

600

601

602

603

604

605

606

607

608

609

610

611

612

613

614

615

616

617

618

619

620

621

622

623

624

625

626

627

628

629

630

631

632

633

634

635

636

637

638

639

640

641

642

643

644

645

646

647

648

649

650

651

652

653

654

655

656

657

658

659

660

661

662

663

664

665

666

667

668

669

670

671

672

673

674

675

676

677

678

679

680

681

682

683

684

685

686

687

688

689

690

691

692

693

694

695

696

697

698

699

700

701

702

703

704

705

706

707

708

709

710

711

712

713

714

715

716

717

718

719

720

721

722

723

724

725

726

727

728

729

730

731

732

733

734

735

736

737

738

739

740

741

742

743

744

745

746

747

748

749

750

751

752

753

754

755

756

757

758

759

760

761

762

763

764

765

766

767

768

769







770

771

```

581 103     output_lines.append("\n    _Esecuzione_test_UDP\n        ...\\n")
582 104     udp_results = run_iperf_test(args.server_ip, 'udp'
583 105         , args.bandwidth, iperf_cmd)
584 106     stats_udp = calculate_stats(udp_results)
585 107     results_summary.append(["UDP", f"{stats_udp['mean']\n        ']:.2f}", f"{stats_udp['max']:.2f}\n        , f"{stats_udp['\n            'stddev']:.2f}"])
586
587
588
589
590 109 # Creazione della tabella con i risultati
591 110 table_header = ["Protocollo", "Media_(Mbps)", "Max_(\n        Mbps)", "Min_(Mbps)", "Dev_Std"]
592 111 table = tabulate(results_summary, headers=\n        table_header, tablefmt="grid")
593
594
595 113 results_output = "\\n    _Risultati:\\n" + table
596 114 print(results_output)
597 115 output_lines.append(results_output)
598
599 117 # Salvataggio dell'output in output.txt
600 118 try:
601 119     with open("output.txt", "w", encoding="utf-8") as
602 120         f:
603 121             f.write("\\n".join(output_lines))
604 122             print("\\n    _Output_salvato_in_output.txt")
605 123 except Exception as e:
606 124     print(f"    _Errore_durante_il_salvataggio_dell'\n            output:{e}")
607
608
609
610
611
612 A.2 Sequence Number
613
614 Sequence Numbers (Distro) for 169.254.27.143:6075 -> 169.254.91.42:5201
615 Ethernet
616
617
618
619
620
621
622
623
624
625 Figure 12: Ethernet to Ethernet sequence number for TCP
626
627 Sequence Numbers (Distro) for 172.21.152.10:6473 -> 172.21.153.12:5201
628 Wi-Fi
629
630
631
632
633
634
635
636
637 Figure 13: Wi-Fi to Wi-Fi sequence number for TCP
638

```

**Figure 14:** ETH to Wi-Fi sequence number for TCP**A.3 More graphs****Figure 15:** Eth to Wi-Fi bits per sec for TCP**Figure 16:** Eth to Wi-Fi bits per sec for UDP**Figure 17:** Wi-Fi to Wi-Fi bits per sec for UDP**Figure 12: Ethernet to Ethernet sequence number for TCP****Figure 13: Wi-Fi to Wi-Fi sequence number for TCP**

639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696